

YOUSSEF SALEM

 youssefkhaledo2@gmail.com  +4917663241140  Berlin, Germany  Youssef Salem  youssefmazenn

PROFESSIONAL EXPERIENCE

AI Research Assistant

University of the Bundeswehr Munich (UniBW) — Technical University of Munich (TUM)

03/2025 – 09/2025
Munich

- Designed **human-centered AI evaluation pipelines** for **LLM-driven conversational agents** in **decision-support scenarios**.
- Architected a **4x2** within-subjects factorial experiment to benchmark multiple **LLMs** under **persona-based prompting**.
- Built a **full-stack experimentation platform** using **Python**, **FastAPI**, and **Next.js** for **scalable data collection**.
- Engineered **LLM persona prompting strategies** manipulating **epistemic certainty** for **comparative behavioral analysis**.
- Performed **statistical analysis** using **repeated-measures ANOVA** and **qualitative thematic analysis**.
- Analyzed **model-level vs. prompt-level effects**, highlighting **model-specific limits of persona control**.
- Conducted research under **cross-institutional supervision (UniBW & TUM)**; presented at TUM and co-authored a submitted paper.

Application specialist

Alliance Healthcare GmbH

09/2024 – 05/2025
Berlin

- Developed **Linux/Unix automation scripts** for **production data workflows** and **job scheduling**, reducing processing time by **35%** and saving **10+ hours per week**.
- Supported and optimized **business-critical data pipelines**, improving **system reliability** and operational efficiency in a **production environment**.
- Re-architected **technical documentation** for automated systems (**50+ pages**), reducing **support tickets by 20%** and improving **new-hire onboarding**.

Database software engineer

Middle East Oil Refinery (MIDOR)

07/2024 – 08/2024

- Built **Python-based ETL automation scripts**, reducing **manual data processing by 10+ hours per week**.
- Implemented **enterprise-scale data validation**, detecting **150+ data inconsistencies** and improving **data integrity**.
- Collaborated with **engineering teams** to translate **business data requirements** into **robust, maintainable pipelines**.

SKILLS

Generative AI & ML: RAG • LangChain • LLM Fine-tuning • Hugging Face Transformers • PyTorch • Scikit-Learn • NLP Data Science & Analytics:
Python • **SQL** • **Pandas** • **NumPy** • **Matplotlib** • **Statistical Modeling** • **EDA Backend & Software Engineering:** FastAPI • REST APIs • Backend System Design • Modular Architectures • Databases (SQL & NoSQL) • Linux / Unix Scripting • Automation **MLOps & Cloud:** Docker • Kubernetes • AWS • CI/CD • Git • Model Packaging & Deployment **AI Systems & Evaluation:** LLM System Design • Persona Prompting • Human-Centered AI Evaluation • Experimental Design • A/B Testing • Statistical Analysis

PROJECTS

Germany RAG Assistant — Production AI System



Technologies: Python, FastAPI, LangChain, ChromaDB, Sentence Transformers, OpenAI API, Docker, Kubernetes, Git, Makefile

- Designed and deployed a **production-grade Retrieval-Augmented Generation (RAG) system** using **LangChain**, **Sentence Transformers**, **ChromaDB**, and **FastAPI**, enabling **citation-grounded LLM responses** over real German government documents and minimizing hallucinations.
- Containerized and orchestrated the **AI backend** with **Docker** and **Kubernetes**, implementing **persistent vector storage (PVCs)**, **Secrets & ConfigMaps**, **health probes**, **resource limits**, and **Horizontal Pod Autoscaling (HPA)** to deliver a **scalable, cloud-ready AI service**.

Conversational AI Agents in VR

Technologies: Unity, C#, FastAPI

- Created **immersive, voice-driven VR NPCs** by integrating **speech-to-text (Whisper)** and **text-to-speech (OpenAI TTS)** pipelines into **Unity** using a **Meta Quest 3** headset.
- Defined a **modular AI dialogue architecture** for **VR NPCs** leveraging locally hosted **LLMs**, achieving **<300 ms latency** for **real-time inference** and enhancing **conversational interactivity**.
- Developed a **Python-based backend solution** using **FastAPI** that processed **200+ requests per second** for **low-latency AI inference**, supporting **continuous, uninterrupted real-time interactions**.

EDUCATION

M.Sc. Computer Science for Digital Media

Berliner Hochschule für Technik

04/2026 – Present
Berlin, Germany

B.Sc. Computer Science and Engineering

German International University in Berlin

10/2021 – 09/2025
Berlin, Germany

LANGUAGES

English

C1, IELTS 7.5

German

C1, Telc C1 Hochschule